

UCI Adult Income Dataset - Exploratory and Descriptive Analysis

In this notebook, we carry out an in-depth exploratory and descriptive analysis of the UCI Adult Income Dataset, a widely used dataset for income prediction tasks based on individual demographic and employment attributes.

This phase of analysis is essential for uncovering patterns, detecting potential biases, and gaining intuition about the dataset's structure before applying any modelling procedures. We examine the distribution of key numerical and categorical variables, investigate relationships between demographic features and income levels, and use visualizations to summarize insights. Particular focus is placed on income disparities across **age groups**, **geographical regions**, **races**, and **education-occupation combinations**, helping lay a solid foundation for downstream modeling and policy-relevant interpretation.

We begin our analysis by importing the core Python libraries required for **data handling**, **numerical computation**, **visualization**, and **directory management**:

- **pandas**: Enables efficient manipulation, filtering, and aggregation of structured tabular data, forming the backbone of our analysis pipeline.
- **numpy**: Provides support for fast numerical operations, array-based computation, and statistical routines.
- **os**: Facilitates interaction with the file system, allowing us to construct flexible and portable directory paths for data and output management.
- **plotly.express**: A high-level graphing library that enables the creation of interactive, publication-quality visualizations, which we use extensively to uncover patterns and present insights throughout the notebook.

```
# import libraries
import pandas as pd
import numpy as np
import os
import plotly.express as px
```

Define and Create Directory Paths

To ensure reproducibility and organized storage, we programmatically create directories if they don't already exist for:

- raw data
- processed data
- results
- documentation

These directories will store intermediate and final outputs for reproducibility.

```
# Get working directory
current_dir = os.getcwd()
# go to directory up to the root directory
project_root_dir = os.path.dirname(current_dir)
# Define paths to the data folders
data_dir = os.path.join(project_root_dir, 'data')
raw_dir = os.path.join(data_dir, 'raw')
processed_dir = os.path.join(data_dir, 'processed')
# Define paths to the data folders
result_dir = os.path.join(project_root_dir, 'result')
# Define paths to docs folder
docs_dir = os.path.join(project_root_dir, 'docs')

# create directories if the do not exist
os.makedirs(raw_dir, exist_ok = True)
os.makedirs(processed_dir, exist_ok = True)
os.makedirs(result_dir, exist_ok = True)
os.makedirs(docs_dir, exist_ok = True)
```

Loading the Cleaned Dataset

We load the cleaned version of the UCI Adult Income Dataset from the processed data directory into a Pandas DataFrame. The `head(10)` function shows the first ten records, giving a glimpse into the data columns such as `age`, `workclass`, `education_num`, etc.

```
adult_data_filename = os.path.join(processed_dir, "adult_cleaned.csv")
adult_df = pd.read_csv(adult_data_filename)
adult_df.head(10)
```

	age	workclass	fnlwgt	education_num	marital_status	relationship	race	sex	
0	39	government	77516	13	single	single	white	male	2
1	50	self-employed	83311	13	married	male spouse	white	male	0
2	38	private	215646	9	divorced or separated	single	white	male	0
3	53	private	234721	7	married	male spouse	black	male	0
4	28	private	338409	13	married	female spouse	black	female	0
5	37	private	284582	14	married	female spouse	white	female	0
6	49	private	160187	5	divorced or separated	single	black	female	0
7	52	self-employed	209642	9	married	male spouse	white	male	0
8	31	private	45781	14	single	single	white	female	1
9	42	private	159449	13	married	male spouse	white	male	5

Dataset Dimensions and Data Types

Here, we examine the structure of the dataset:

- There are *32,513* entries and *16* variables.
- The dataset includes both **numerical** (e.g., `age`, `hours_per_week`) and **categorical** variables (e.g., `sex`, `education_level`).

Understanding data types and null entries is essential before proceeding with analysis.

```
adult_df.shape
```

```
(32533, 16)
```

```
adult_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32533 entries, 0 to 32532
Data columns (total 16 columns):
#   Column              Non-Null Count  Dtype
---  -
0   age                 32533 non-null  int64
1   workclass           32533 non-null  object
2   fnlwgt              32533 non-null  int64
3   education_num       32533 non-null  int64
4   marital_status      32533 non-null  object
5   relationship        32533 non-null  object
6   race                32533 non-null  object
```

```

7  sex                32533 non-null  object
8  capital_gain       32533 non-null  int64
9  capital_loss       32533 non-null  int64
10 hours_per_week     32533 non-null  int64
11 income             32533 non-null  object
12 education_level    32533 non-null  object
13 occupation-grouped 32533 non-null  object
14 native_region      32533 non-null  object
15 age_group          32533 non-null  object
dtypes: int64(6), object(10)
memory usage: 4.0+ MB

```

summary statistics

Numerical variables

```
adult_df.describe() # summaries the numerical data
```

	age	fnlwgt	education_num	capital_gain	capital_loss	hours_per_week
count	32533.000000	3.253300e+04	32533.000000	32533.000000	32533.000000	32533.000000
mean	38.587557	1.897849e+05	10.081640	1078.576338	87.378969	40.441306
std	13.637609	1.055601e+05	2.571689	7388.401928	403.125450	12.346494
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178330e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.369930e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

This summary provides a snapshot of key distribution characteristics. We see that:

- Age ranges from 17 to 90, with a mean of 38.6 years. It is slightly right-skewed (positively skewed). While the average age is approximately 38.6 years, an examination of the percentiles reveals that the majority of individuals are clustered in the younger to middle-age range, with fewer observations in the older age brackets. This skewed age distribution might suggest labor force participation is concentrated in specific age groups, which could reflect broader demographic or economic realities.
- Capital gains/losses are highly skewed, with most values at 0 (the 75th percentile is 0). This indicates that a small number of individuals report very large gains or losses,

especially evident in the capital gain variable which reaches up to \$99,999. These variables act as proxies for wealth-related income that goes beyond regular wages or salaries. Individuals with non-zero values for capital gains or losses often represent a distinct socioeconomic subset of the population — typically more financially literate, or with access to investment assets. The stark inequality in their distributions mirrors real-world disparities in asset ownership and investment returns.

- The dataset has individuals working anywhere from 1 to 99 hours per week, with a median of 40. This aligns with the standard full-time work week in many countries (8 hours per day for 5 working days). The mean is slightly above that at 40.4 hours, suggesting a mild right skew, with a small subset of individuals working significantly longer hours. The mode is also 40, further reinforcing the prevalence of full-time work. A non-trivial number of individuals report working very few hours, possibly due to part-time work, unemployment, or semi-retirement. On the other extreme, some report working more than 45 hours per week, which may indicate multiple jobs, weekend-work, self-employment, or informal labor, and could reflect socioeconomic necessity.

Categorical variables

```
adult_df.describe(include='object')
```

	workclass	marital_status	relationship	race	sex	income	education_level	occupation
count	32533	32533	32533	32533	32533	32533	32533	32533
unique	6	4	5	5	2	2	7	9
top	private	married	male spouse	white	male	<=50k	secondary graduate	white coll
freq	22670	14993	13187	27791	21774	24694	10494	11480

```
adult_df['workclass'].value_counts(normalize=True)
```

```
workclass
private      0.696831
government   0.133710
self-employed 0.112378
unknown      0.056435
voluntary    0.000430
unemployed   0.000215
Name: proportion, dtype: float64
```

```
adult_df['marital_status'].value_counts(normalize=True)
```

```
marital_status
married          0.460855
single           0.327760
divorced or separated  0.180863
widowed          0.030523
Name: proportion, dtype: float64
```

```
adult_df['relationship'].value_counts(normalize=True)
```

```
relationship
male spouse    0.405342
single         0.360680
child          0.155627
female spouse  0.048197
extended relative  0.030154
Name: proportion, dtype: float64
```

```
adult_df['race'].value_counts(normalize=True)
```

```
race
white          0.854240
black          0.095964
asian or pacific islander  0.031906
american indian or eskimo  0.009560
other          0.008330
Name: proportion, dtype: float64
```

```
adult_df['sex'].value_counts(normalize=True)
```

```
sex
male    0.66929
female  0.33071
Name: proportion, dtype: float64
```

```
adult_df['education_level'].value_counts(normalize=True)
```

```

education_level
secondary graduate    0.322565
tertiary              0.247810
unemployed            0.223773
secondary             0.093905
associate             0.075277
primary               0.035134
preschool             0.001537
Name: proportion, dtype: float64

```

```
adult_df['occupation-grouped'].value_counts(normalize=True)
```

```

occupation-grouped
white collar    0.352872
blue collar    0.156334
white color    0.127132
service        0.125626
blue collar    0.091169
whitecollar    0.061476
unknown        0.056650
white collar    0.028463
military        0.000277
Name: proportion, dtype: float64

```

```
adult_df['native_region'].value_counts(normalize=True)
```

```

native_region
north america    0.923278
asia             0.020625
other            0.017890
central america  0.016107
europe           0.016015
south america    0.006086
Name: proportion, dtype: float64

```

```
adult_df['age_group'].value_counts(normalize=True)
```

```

age_group
26-35    0.261581
36-45    0.246058

```

```
46-60    0.224111
18-25    0.167553
61-75    0.064273
<18      0.029047
76+      0.007377
Name: proportion, dtype: float64
```

`workclass`

The private sector dominates, employing ~69.7% of the population. The government sector (13.4%) and self-employment (11.2%) also make up substantial portions of the workforce. A small fraction is labeled as “unknown” (5.6%), which may correspond to missing or ambiguous data entries. Tiny proportions are voluntary (0.04%) or unemployed (0.02%), possibly underreported or underrepresented in the sample.

`marital_status`

Married individuals make up the largest group (46.1%), followed by those who are single (32.8%) and divorced or separated (18.1%). Widowed individuals represent a small minority (~3.1%).

`relationship`

The majority are labeled as “male spouse” (40.5%) or “single” (36.1%). Smaller categories include children (15.6%), female spouses (4.8%), and extended relatives (3.0%). The dominance of `male spouse` reflects the dataset’s gendered structure and may point to traditional family roles. The relative scarcity of “female spouse” roles suggests potential gender imbalances in how income-earning is reported within households.

`race`

The dataset is overwhelmingly composed of White individuals (~85.4%). Other racial groups include Black (9.6%), Asian or Pacific Islander (3.2%), American Indian or Eskimo (1.0%), and Other (0.8%). The racial imbalance limits the generalizability of models trained on this data. Smaller racial groups may suffer from limited statistical power, affecting fairness and performance in predictive modeling.

`sex`

Males constitute 66.9% of the dataset, with females making up the remaining 33.1%. This male-skewed distribution could be due to sampling (e.g., primary earners in households), workforce participation patterns, or reporting biases.

`education_level`

Secondary-school graduates form the largest educational group (~32%), highlighting the central role of high school completion in the labor force. Tertiary education holders — those with university or equivalent degrees — account for nearly 25% of the population, representing a

substantial segment with advanced qualifications. A notable 22.4% have attended some college without necessarily earning a degree, suggesting that partial post-secondary education is common, yet may not always translate into formal certification. The remaining 20% are distributed among those with only secondary education (9.4%), associate degrees (7.5%), primary school (3.5%), and a very small group with only preschool education (0.15%). It is evident that the education distribution is skewed toward mid- to high-level education, with relatively few individuals having only basic schooling. This reflects a dataset that largely captures working-age adults in formal labor, which may underrepresent the least-educated populations.

`occupation_grouped`

White-collar occupations are the most prevalent (~51%), followed by blue-collar, service, and unknown. Smaller categories include military, which is marginal. Essentially, slightly over half of individuals in the dataset work in professional, managerial, sales, clerical, or tech-support roles. This suggests the dataset is heavily weighted toward professional and administrative occupations. Nearly a third of the population works in manual labor or skilled trade positions (craft, transport, machine operation, farming, etc.). This indicates a significant segment engaged in physically intensive or technical labor.

`native_region`

The vast majority of individuals are from North America (~92.3%). Smaller proportions are from Central America, Asia, Europe, South America, and a generic Other category. The heavy concentration of North American individuals reflects the U.S. focus of the dataset.

`age_group`

The largest groups are 26–35 and 36–45, followed by 46–60. These three age groups represent about 73% of the dataset. Very few individuals are under 18 or above 75, consistent with the dataset's focus on the working-age population.


Income Distribution

Given that `income` is the target variable, most of the analysis hereafter will be based on it. We first of all examine the income distribution in the dataset.

```
adult_df_income = adult_df.groupby('income').size().reset_index(name='total')
adult_df_income
```

	income	total
0	<=50k	24694
1	>50k	7839


```
fig = px.pie(adult_df_income, names='income' , values='total', title='Overall Income Distribution')
fig.update_layout(
    template='plotly_dark', # Safe background styling
    paper_bgcolor='black'
)
fig.show()
fig.write_image(os.path.join(result_dir, 'income_distribution_pie_chart.jpg'))
fig.write_image(os.path.join(result_dir, 'income_distribution_pie_chart.png'))
fig.write_html(os.path.join(result_dir, 'income_distribution_pie_chart.html'))
```



WARNING Thread(Thread-57 (run)) Task(Task-944) choreographer.browser_async:browser_async.py:.

```
fig = px.pie(adult_df_income, names='income' , values='total', title='Overall Income Distribution')
fig.show()
```

Overall Income Distribution



This pie chart visualizes the overall income split: 76% of individuals earn ≤50K, while 24% earn >50K. This means that nearly 3 out of 4 individuals fall into the lower income bracket (≤50K). This shows that there is a significant imbalance.

Income by Age group

```
adult_df_income_age= adult_df.groupby(['age_group', 'income']).size().reset_index(name='total_by_age')
adult_df_income_age
```

	age_group	income	total_by_age
0	18-25	<=50k	5337
1	18-25	>50k	114
2	26-35	<=50k	6919
3	26-35	>50k	1591
4	36-45	<=50k	5233
5	36-45	>50k	2772
6	46-60	<=50k	4480
7	46-60	>50k	2811
8	61-75	<=50k	1580
9	61-75	>50k	511
10	76+	<=50k	200
11	76+	>50k	40
12	<18	<=50k	945

```
adult_df_income_age = adult_df.groupby(['age_group', 'income']).size().reset_index(name='total_by_age')
adult_df_income_age
```

	age_group	income	total_by_age
0	18-25	<=50k	5337
1	18-25	>50k	114
2	26-35	<=50k	6919
3	26-35	>50k	1591
4	36-45	<=50k	5233
5	36-45	>50k	2772
6	46-60	<=50k	4480
7	46-60	>50k	2811
8	61-75	<=50k	1580
9	61-75	>50k	511
10	76+	<=50k	200
11	76+	>50k	40
12	<18	<=50k	945

```
total_per_group = adult_df_income_age.groupby('age_group').size()
total_per_group
```

```
age_group
18-25    2
26-35    2
36-45    2
46-60    2
61-75    2
76+      2
<18      1
dtype: int64
```

```
total_per_group = adult_df_income_age.groupby('age_group')['total_by_age'].transform('sum')
total_per_group
```


```
0    5451
1    5451
2    8510
3    8510
4    8005
5    8005
6    7291
7    7291
8    2091
9    2091
10   240
11   240
12   945
Name: total_by_age, dtype: int64
```

```
total_per_group = adult_df_income_age.groupby('age_group')['total_by_age'].transform('sum')
adult_df_income_age['percentage'] = (adult_df_income_age['total_by_age']/total_per_group) * 100
adult_df_income_age
```

	age_group	income	total_by_age	percentage
0	18-25	<=50k	5337	97.908641
1	18-25	>50k	114	2.091359
2	26-35	<=50k	6919	81.304348

	age_group	income	total_by_age	percentage
3	26-35	>50k	1591	18.695652
4	36-45	<=50k	5233	65.371643
5	36-45	>50k	2772	34.628357
6	46-60	<=50k	4480	61.445618
7	46-60	>50k	2811	38.554382
8	61-75	<=50k	1580	75.561932
9	61-75	>50k	511	24.438068
10	76+	<=50k	200	83.333333
11	76+	>50k	40	16.666667
12	<18	<=50k	945	100.000000

```
fig = px.bar(
    adult_df_income_age,
    x = 'age_group',
    y = 'percentage',
    color = 'income',
    title='Income Distribution by Age Group(%)',
    barmode='group',
    color_discrete_sequence=['#0d2e03', '#56b33a'],
)
fig.update_traces(texttemplate = 'Updated %')
fig.show()
fig.write_image(os.path.join(result_dir, 'income_distribution_by_agegroupchart.jpg'))
fig.write_image(os.path.join(result_dir, 'income_distribution_pie_chart.png'))
fig.write_image(os.path.join(result_dir, 'income_distribution_pie_chart.html'))
themes = ["plotly", "plotly_white", "plotly_dark", "ggplot2", "seaborn", "simple_white", "pres
for theme in themes:
    fig.update_layout(template=theme)
    fig.show()
```



WARNING Thread(Thread-65 (run)) Task(Task-1080) choreographer.browser_async:browser_async.py

RuntimeError: Couldn't close or kill browser subprocess

RuntimeError Traceback (most recent call last)

Cell In[75], line 12

```
10 fig.update_traces(texttemplate = 'Updated %')
11 fig.show()
--> 12 fig.write_image(os.path.join(result_dir, 'income_distribution_by_agegroupchart.jpg'))
13 fig.write_image(os.path.join(result_dir, 'income_distribution_pie_chart.png'))
14 fig.write_image(os.path.join(result_dir, 'income_distribution_pie_chart.html'))
```

File ~\anaconda3\Lib\site-packages\plotly\basedatatypes.py:3911, in BaseFigure.write_image(s

```
3907     if kwargs.get("engine", None):
3908         warnings.warn(
3909             ENGINE_PARAM_DEPRECATION_MSG, DeprecationWarning, stacklevel=2
3910         )
```

-> 3911 return pio.write_image(self, *args, **kwargs)

File ~\anaconda3\Lib\site-packages\plotly\io_kaleido.py:509, in write_image(fig, file, form

```
505 format = infer_format(path, format)
507 # Request image
508 # Do this first so we don't create a file if image conversion fails
--> 509 img_data = to_image(
```

```
510     fig,
511     format=format,
512     scale=scale,
513     width=width,
514     height=height,
515     validate=validate,
516     engine=engine,
517 )
519 # Open file
520 if path is None:
521     # We previously failed to make sense of `file` as a pathlib object.
522     # Attempt to write to `file` as an open file descriptor.
```

File ~\anaconda3\Lib\site-packages\plotly\io_kaleido.py:373, in to_image(fig, format, width

```
369 from kaleido.errors import ChromeNotFoundError
371 try:
372     # TODO: Refactor to make it possible to use a shared Kaleido instance here
--> 373     img_bytes = kaleido.calc_fig_sync(
374         fig_dict,
375         opts=dict(
376             format=format or defaults.default_format,
377             width=width or defaults.default_width,
```

```

378         height=height or defaults.default_height,
379         scale=scale or defaults.default_scale,
380     ),
381     topojson=defaults.topojson,
382     kopts=(
383         dict(
384             mathjax=defaults.mathjax,
385         )
386         if defaults.mathjax
387         else None
388     ),
389 )
390 except ChromeNotFoundError:
391     raise RuntimeError(PLOTLY_GET_CHROME_ERROR_MSG)
File ~\anaconda3\Lib\site-packages\kaleido\__init__.py:145, in calc_fig_sync(*args, **kwargs)
143 def calc_fig_sync(*args, **kwargs):
144     """Call `calc_fig` but blocking."""
--> 145     return _async_thread_run(calc_fig, args=args, kwargs=kwargs)
File ~\anaconda3\Lib\site-packages\kaleido\__init__.py:138, in _async_thread_run(func, args,
136 res = q.get()
137 if isinstance(res, BaseException):
--> 138     raise res
139 else:
140     return res
File ~\anaconda3\Lib\site-packages\kaleido\__init__.py:129, in _async_thread_run.<locals>.run
126 def run(*args, **kwargs):
127     # func is a closure
128     try:
--> 129         q.put(asyncio.run(func(*args, **kwargs)))
130     except BaseException as e: # noqa: BLE001
131         q.put(e)
File ~\anaconda3\Lib\asyncio\runners.py:194, in run(main, debug, loop_factory)
190     raise RuntimeError(
191         "asyncio.run() cannot be called from a running event loop")
193 with Runner(debug=debug, loop_factory=loop_factory) as runner:
--> 194     return runner.run(main)
File ~\anaconda3\Lib\asyncio\runners.py:118, in Runner.run(self, coro, context)
116 self._interrupt_count = 0
117 try:
--> 118     return self._loop.run_until_complete(task)
119 except exceptions.CancelledError:
120     if self._interrupt_count > 0:
File ~\anaconda3\Lib\asyncio\base_events.py:687, in BaseEventLoop.run_until_complete(self, f

```

```

684 if not future.done():
685     raise RuntimeError('Event loop stopped before Future completed.')
--> 687 return future.result()
File ~\anaconda3\Lib\site-packages\kaleido\__init__.py:54, in calc_fig(fig, path, opts, topo
52 kopts = kopts or {}
53 kopts["n"] = 1
--> 54 async with Kaleido(**kopts) as k:
55     return await k.calc_fig(
56         fig,
57         path=path,
58         opts=opts,
59         topojson=topojson,
60     )
File ~\anaconda3\Lib\site-packages\kaleido\kaleido.py:76, in Kaleido.__aexit__(self, exc_type
74 await asyncio.gather(*self._background_render_tasks, return_exceptions=True)
75 _logger.info("Exiting Kaleido")
--> 76 return await super().__aexit__(exc_type, exc_value, exc_tb)
File ~\anaconda3\Lib\site-packages\choreographer\browser_async.py:249, in Browser.__aexit__(
242 async def __aexit__(
243     self,
244     type_: type[BaseException] | None,
245     value: BaseException | None,
246     traceback: TracebackType | None,
247 ) -> None: # None instead of False is fine, eases type checking
248     """Close the browser."""
--> 249     await self.close()
File ~\anaconda3\Lib\site-packages\kaleido\kaleido.py:69, in Kaleido.close(self)
67     task.cancel()
68 _logger.info("Exiting Kaleido/Choreo")
--> 69 return await super().close()
File ~\anaconda3\Lib\site-packages\choreographer\browser_async.py:228, in Browser.close(self)
226 try:
227     _logger.debug("Starting browser close methods.")
--> 228     await self._close()
229     _logger.debug("Browser close methods finished.")
230 except ProcessLookupError:
File ~\anaconda3\Lib\site-packages\choreographer\browser_async.py:216, in Browser._close(self)
214     return
215 else:
--> 216     raise RuntimeError("Couldn't close or kill browser subprocess")
RuntimeError: Couldn't close or kill browser subprocess

```



```
pip install -U kaleido
```

```
pip install -U plotly
```

```
adult_df_income_native_region = adult_df.groupby(['native_region', 'income']).size().reset_index()
adult_df_income_native_region
```

	native_region	income	total_income_distr
0	asia	<=50k	465
1	asia	>50k	206
2	central america	<=50k	466
3	central america	>50k	58
4	europe	<=50k	369
5	europe	>50k	152
6	north america	<=50k	22784
7	north america	>50k	7253
8	other	<=50k	436
9	other	>50k	146
10	south america	<=50k	174
11	south america	>50k	24


```
total_per_region = adult_df_income_native_region.groupby('native_region')['total_income_distr'].sum()
adult_df_income_native_region['percentage'] = (adult_df_income_native_region['total_income_distr'] / total_per_region) * 100
adult_df_income_native_region
```

	native_region	income	total_income_distr	percentage
0	asia	<=50k	465	69.299553
1	asia	>50k	206	30.700447
2	central america	<=50k	466	88.931298
3	central america	>50k	58	11.068702
4	europe	<=50k	369	70.825336
5	europe	>50k	152	29.174664
6	north america	<=50k	22784	75.853114
7	north america	>50k	7253	24.146886
8	other	<=50k	436	74.914089
9	other	>50k	146	25.085911
10	south america	<=50k	174	87.878788
11	south america	>50k	24	12.121212

```

fig = px.bar(
    adult_df_income_native_region,
    x = 'native_region',
    y = 'percentage',
    color = 'income',
    title='Income Distribution by native_region(%)',
    barmode='group',
    color_discrete_sequence=['#0d2e03', '#56b33a'],
    text='percentage'
)
fig.update_traces(texttemplate = '%{text:.2f}%')
fig.update_layout(template='presentation', paper_bgcolor='rgb(0, 0, 0)', plot_bgcolor='rgb(0, 0, 0)')
fig.show()
fig.write_image(os.path.join(result_dir, 'income_distribution_by_nativeregion_bar_plot.jpg'))
fig.write_image(os.path.join(result_dir, 'income_distribution_by_nativeregion_bar_plot.png'))
fig.write_html(os.path.join(result_dir, 'income_distribution_by_nativeregion_bar_plot.html'))

```




WARNING Thread(Thread-71 (run)) Task(Task-1199) choreographer.browser_async:browser_async.py

Asia (30.7%) and Europe (29.2%) have the highest proportions of high-income earners. This suggests these immigrant groups might be better integrated into high-paying professional roles, or may represent a more skilled migrant profile in the dataset. Central America (11.1%) and South America (12.1%) have the lowest proportions of >50K earners. With 24.2% of North Americans earning >50K, this serves as a middle-ground baseline. Interestingly, both Asian and European groups outperform the native-born population proportionally in high-income brackets. The 'Other' group sits around 25.1%, close to North America's rate. This likely reflects a diverse mix of regions not explicitly listed.

```

fig = px.bar(
    adult_df_income_age,
    x = 'age_group',
    y = 'percentage',
    color = 'income',
    title='Income Distribution by Age Group(%)',
    barmode='group',
    color_discrete_sequence=['#0d2e03', '#56b33a'],
    text='percentage'
)
fig.update_traces(texttemplate = '%{text:.2f}%')
fig.update_layout(template='presentation', paper_bgcolor='rgb(0, 0, 0)', plot_bgcolor='rgb(0, 0, 0)')
fig.show()
fig.write_image(os.path.join(result_dir, 'income_distribution_by_agegroup_bar_plot.jpg'))
fig.write_image(os.path.join(result_dir, 'income_distribution_by_agegroup_bar_plot.png'))
fig.write_html(os.path.join(result_dir, 'income_distribution_by_agegroup_bar_plot.html'))

```



WARNING Thread(Thread-75 (run)) Task(Task-1275) choreographer.browser_async:browser_async.py

```

# Group by native_region and income, count occurrences
adult_df_income_reg = adult_df.groupby(['native_region', 'income']).size().reset_index(name='count')

# Calculate total per native_region
total_per_group = adult_df_income_reg.groupby('native_region')['count'].transform('sum')

# Calculate percentage
adult_df_income_reg['percentage'] = (adult_df_income_reg['count'] / total_per_group) * 100

# Plot the bar chart


```

```

fig = px.bar(
    adult_df_income_reg,
    x='native_region',
    y='percentage',
    color='income',
    title='Income Distribution by Native Region (%)',
    barmode='group',
    color_discrete_sequence=['#0d2e03', '#56b33a'],
    text='percentage'
)

# Format the text on bars
fig.update_traces(texttemplate='%{text:.2f}%', textposition='outside')
fig.update_layout(template='presentation', paper_bgcolor='rgb(0, 0, 0)', plot_bgcolor='rgb(0, 0, 0)')
fig.show()
fig.write_image(os.path.join(result_dir, 'income_distribution_by_race_bar_plot.jpg'))
fig.write_image(os.path.join(result_dir, 'income_distribution_by_race_bar_plot.png'))
fig.write_html(os.path.join(result_dir, 'income_distribution_by_race_bar_plot.html'))

```



```

# Group by native_region and income, count occurrences
adult_df_income_reg = adult_df.groupby(['native_region', 'income']).size().reset_index(name='count')

# Calculate total per native_region
total_per_group = adult_df_income_reg.groupby('native_region')['count'].transform('sum')

# Calculate percentage
adult_df_income_reg['percentage'] = (adult_df_income_reg['count'] / total_per_group) * 100

# Plot the bar chart
fig = px.bar(


```

```

adult_df_income_reg,
x='native_region',
y='percentage',
color='income',
title='Income Distribution by Native Region (%)',
barmode='group',
color_discrete_sequence=['#0d2e03', '#56b33a'],
text='percentage'
)

# Format the text on bars
fig.update_traces(texttemplate='%{text:.2f}%', textposition='outside')
fig.update_layout(template='presentation', paper_bgcolor='rgb(0, 0, 0)', plot_bgcolor='rgb(0, 0, 0)')
fig.show()
fig.write_image(os.path.join(result_dir, 'income_distribution_by_race_bar_plot.jpg'))
fig.write_image(os.path.join(result_dir, 'income_distribution_by_race_bar_plot.png'))
fig.write_html(os.path.join(result_dir, 'income_distribution_by_race_bar_plot.html'))

```



```

# Group by native_region and income, count occurrences
adult_df_income_reg = adult_df.groupby(['race', 'income']).size().reset_index(name='total_income_reg')

# Calculate total per native_region
total_per_group = adult_df_income_reg.groupby('race')['total_income_reg'].transform('sum')

# Calculate percentage
adult_df_income_reg['percentage'] = (adult_df_income_reg['total_income_reg'] / total_per_group) * 100

# Plot the bar chart
fig = px.bar(
    adult_df_income_reg,

```

```

x='race',
y='percentage',
color='income',
title='Income Distribution by Race (%)',
barmode='group',
color_discrete_sequence=['#0d2e03', '#56b33a'],
text='percentage'
)

# Format the text on bars
fig.update_traces(texttemplate='%{text:.2f}%', textposition='outside')
fig.update_layout(template='presentation', paper_bgcolor='rgb(0, 0, 0)', plot_bgcolor='rgb(0, 0, 0)')
fig.show()
fig.write_image(os.path.join(result_dir, 'income_distribution_by_race_bar_plot.jpg'))
fig.write_image(os.path.join(result_dir, 'income_distribution_by_race_bar_plot.png'))
fig.write_html(os.path.join(result_dir, 'income_distribution_by_race_bar_plot.html'))

```



WARNING Thread(Thread-83 (run)) Task(Task-1435) choreographer.browser_async:browser_async.py

```

# Group by native_region and income, count occurrences
adult_df_income_reg = adult_df.groupby(['race', 'income']).size().reset_index(name='total_income_reg')

# Calculate total per native_region
total_per_group = adult_df_income_reg.groupby('race')['total_income_reg'].transform('sum')

# Calculate percentage
adult_df_income_reg['percentage'] = (adult_df_income_reg['total_income_reg'] / total_per_group)

# Plot the bar chart

```

```

fig = px.bar(
    adult_df_income_reg,
    x='race',
    y='percentage',
    color='income',
    title='Income Distribution by Race (%)',
    barmode='group',
    color_discrete_sequence=['#0d2e03', '#56b33a'],
    text='percentage'
)

# Format the text on bars
fig.update_traces(texttemplate='%{text:.2f}%', textposition='outside')
fig.update_layout(template='presentation', paper_bgcolor='rgb(0, 0, 0)', plot_bgcolor='rgb(0, 0, 0)')
fig.show()
fig.write_image(os.path.join(result_dir, 'income_distribution_by_race_bar_plot.jpg'))
fig.write_image(os.path.join(result_dir, 'income_distribution_by_race_bar_plot.png'))
fig.write_html(os.path.join(result_dir, 'income_distribution_by_race_bar_plot.html'))

```



WARNING Thread(Thread-83 (run)) Task(Task-1435) choreographer.browser_async:browser_async.py

```

adult_df_income_edu_occ = adult_df.groupby(['education_level', 'income', 'occupation-grouped'])
adult_df_income_edu_occ

```

	education_level	income	occupation-grouped	total
60	secondary graduate	<=50k	white collar	2617
95	unemployed	<=50k	white collar	2497
54	secondary graduate	<=50k	blue collar	1871

	education_level	income	occupation-grouped	total
86	tertiary	>50k	white collar	1865
77	tertiary	<=50k	white collar	1667
...
39	secondary	<=50k	military	1
33	primary	>50k	service	1
21	preschool	<=50k	white color	1
82	tertiary	>50k	military	1
73	tertiary	<=50k	military	1

```
final_file = os.path.join(processed_dir, 'adult_cleaned.csv')
adult_df.to_csv(final_file, index=False)
```